

QRM II Graded Assignment (6)

Period 1, 2026

group 6,

26-09-2026

0 Introduction

This assignment is to be completed in groups of 4-5. Further, all students in your group need to be assigned to the same R tutorial group (Friday's tutorial). You can sign yourself up for a group on Canvas. Please do so **before the start of your first R tutorial on Friday September 4th**. You can use the Discussion Board in Canvas if you do not have a group yet or if your group is incomplete.

The assignment has 5 parts, and each part corresponds to the course material of that week (with the exclusion of week 6, for which there is no R programming material).

You are supposed to hand in these assignments on Canvas at the following dates:

- **Deadline 1** *Sunday September 27th, at 23:59pm*: you are supposed to hand in weeks 1, 2, and 3 of this assignment. This will determine 18% of your overall course grade
- **Deadline 2** *Sunday October 11th, at 23:59pm*: you are supposed to hand in weeks 4, and 5 of this assignment. This will determine 12% of your overall course grade

The R tutorials (each Friday) will consist of two halves. During the first half, you will discuss the tutorial exercises. These can be downloaded separately from Canvas. During the second half, you can work on this graded assignment within your own group. The purpose is that you find out how to work with R for doing statistical analyses by yourself. The tutorial exercises are meant to teach you basic commands to get you started, but to answer the problem sets in this assignment, you might need to research your own solutions, and use functions and commands not described in the tutorial exercises. Learning how to solve your own research problems is integral part of learning R. When you and your group get stuck on how to approach an exercise, the hierarchy in finding your way is as follows:

- use the concepts from the tutorial exercises;
- use the cheat sheets available on Canvas;
- use Google, YouTube, StackOverflow, or another website;
- ask the teacher.

The use of generative AI is **not** permitted and may result in a grade of 0. See the AI protocol in the course manual for details.

To answer the assignment, you can simply fill out this R markdown document. There are designated places which you can fill with R code. There are also designated spaces for you to answer each question. Often, the structure of an answer will be as follows. First, you type the R code in the designated box. This will show how you analyzed the data to get the answer to the question. Below the box for the R code, you will then summarize your answer to the question, i.e. what are the conclusions that you draw from the data analysis?

When handing in, you are supposed to submit this .Rmd file, and a knitted version of this document. You can knit this document to pdf, word, or html. Knitting to pdf requires you to have a .tex distribution installed on your computer. Knitting to Word requires you to have Word installed.

The exercises are designed such that you should be able to finish the majority of them during the tutorial each week. If you are not able to finish them fully during that time, you are expected to work on it in your own time using the computers on campus or your own device. It is best to meet as a group in-person when working together. If you want to work remotely, github is a good platform to guarantee smooth collaboration. Alternatively, you can email this .Rmd file back and forth to one another as a group, but this is not recommended as it is more cumbersome.

We encourage you to keep your code blocks, printing statements, and final answers, as short as possible. In any case, there is a page limit of 6 pages per week, which encompasses the total length of this document which consists of the questions, your coding lines, and your answers. When your answers to questions of the respective week exceed this page limit, they will not be graded, resulting in zero points.

1 Week 1

1. Find the dataset “movies4.tsv” on Canvas. Describe your data set: How many observations does it have. How many variables are there? How many subjects? What consists of a subject? **[4 points]**

```
#WRITE YOUR CODE HERE
movies4 <- read.delim("movies4.tsv", header = TRUE)
nrow(movies4)
```

```
## [1] 808
```

```
ncol(movies4)
```

```
## [1] 19
```

```
#hernoemen row naar subjects en column naar variables keer elkaar = observations
subjects <- nrow(movies4)
variables <- ncol(movies4)
observations <- subjects * variables
print(subjects)
```

```
## [1] 808
```

```
print(variables)
```

```
## [1] 19
```

```
print(observations)
```

```
## [1] 15352
```

Your Answer:

The dataset has 15352 observations, 19 variables, 808 subjects. A subject consists of a unit that is being researched. In this case different movies.

2. Which of the following types of variables are present in your data set? (i) nominal; (ii) ordinal; (iii); continuous; (iv) discrete. If present, name one example of such a variable present in your data set. [4 points]

```
head(movies4, n=4)
```

```
##   index  budget                                keywords
## 1  2256 0.0e+00                                <NA>
## 2  2473 1.5e+07          barcelona spain menage a trois author
## 3  1514 3.2e+07          gunslinger revenge prairie shootout pistol
## 4  1512 3.2e+07 robbery double life dual identity small town indiana
##   original_language                                title popularity
## 1                  en The Oogieloves in the Big Balloon Adventure    0.28576
## 2                  en              Vicky Cristina Barcelona    32.75825
## 3                  en              The Quick and the Dead    16.47336
## 4                  en              A History of Violence    34.62874
##   release_date  revenue runtime vote_average vote_count  genre release_year
## 1  2012-08-29         0      88         2.2         8 Family         2012
## 2  2008-08-15 96408652      96         6.7        1020 Drama         2008
## 3  1995-02-09 18552460     107         6.3         427 Action         1995
## 4  2005-09-23 60740827      96         6.9         832 Drama         2005
##   release_month release_day  first_actor first_actor_gender
## 1              8          29 Christopher Lloyd           male
## 2              8          15 Scarlett Johansson          female
## 3              2           9   Sharon Stone            female
## 4              9          23   Viggo Mortensen            male
##   director_first_name director_gender
## 1             Matthew           male
## 2              Woody           male
## 3               Sam            male
## 4              David            male
```

Your Answer:

Genre is nominal, ordinal is not present, popularity is continuous and runtime is discrete.

3. A movie studio wants to know which types of movies give maximal revenue. Perform the following steps to provide the movie studio with an analysis which corresponds to their request:
- Summarize the variable revenue and report its mean, median, maximum, and minimum. [2 points]
 - Which movie has the highest revenue in your data set and how much is this revenue. Which movie has the lowest and how much is its revenue? If multiple movies have the exact same highest or lowest revenue, give only one example. [2 points]
 - Create a histogram of the variable revenue. Make sure it has an appropriate title, and appropriate titles and labels for the x- and y-axis. Describe the shape of the histogram. What does this tell you about the nature of making money in the movies industry? [2 points]
 - Add a new variable to your data set: the log of revenue. When creating this variable, what happens to movies for which revenue is zero? What then happens when you calculate the mean of log revenue? [2 points]

- e. For movies that have a revenue of zero, replace log of revenue with “NA”. What is now the mean of log of revenue? Create a histogram for log revenue, again with an appropriate title, x- and y-axis labels. Describe the shape of this histogram. **[2 points]**
- f. For the variables revenue and log of revenue calculate their skewness and kurtosis. How do these numbers relate to the shape of the histograms that you estimated under d.) and e.)? **[2 points]**

For each step, you should provide first all the code you used to answer the question and then formulate an answer using full sentences.

Each week consists of 1, 2, or 3 subquestions. The total amount of points you can earn per week is 20 points.

Step a

#WRITE YOUR CODE HERE

```
movies4 %>%
  summarise(mean_revenue = mean(revenue),
            median_revenue = median(revenue),
            maximum_revenue = max(revenue),
            minimum_revenue = min(revenue))
```

```
##   mean_revenue median_revenue maximum_revenue minimum_revenue
## 1      82727885      14560504      1519557910              0
```

Your Answer:

The mean is 82727885, median is 14560504, maximum is 1519557910 and the minimum is 0. *Step b*

```
max(movies4$revenue)
```

```
## [1] 1519557910
```

```
min(movies4$revenue)
```


```
## [1] 0
```

Your Answer:

the movie with the highest revenue is The Avengers with a revenue of 1519557910 and an example of a movie with the lowest revenue is Red Dog with a revenue of 0. *Step c*

```
hist(movies4$revenue,
     main = "Revenue of movies",
     xlab = "Revenue",
     ylab = "Frequency")
```

Revenue of movies



Your Answer:

The histogram is skewed right. This tells us that it is hard to make Revenue from movies in the movie industry. Most movies will not earn a lot. *Step d*

```
movies_LogRevenue <- log(movies4$revenue)
print(movies_LogRevenue)
```

```
## [1]      -Inf 18.384107 16.736113 17.922127 19.084238      -Inf      -Inf
## [8]      -Inf      -Inf 10.747100 18.918470 18.637984      -Inf 18.255485
## [15] 18.490369      -Inf      -Inf      -Inf 18.372660 19.225658      -Inf
## [22] 20.105195 19.125227 17.196247      -Inf 18.972547      -Inf      -Inf
## [29] 16.578277      -Inf      -Inf 18.293515 17.801258 17.031271 19.052339
## [36] 12.706993      -Inf      -Inf      -Inf 18.543078 15.407136      -Inf
## [43] 19.346038 17.053133      -Inf 18.264990 20.603002 16.119826 14.900016
## [50]      -Inf 16.090411 15.087134 19.692163 18.877394 19.654616 14.976503
## [57] 17.825704 14.945238 20.865897 12.945533 10.947732 12.726270 20.590829
## [64] 18.095755 17.455866 17.067303 14.856785      -Inf 17.078275 19.176213
## [71] 19.153415      -Inf 18.263396      -Inf 15.908505 15.158411      -Inf
## [78]      -Inf 17.278518 3.135494 16.523561      -Inf 17.372506 16.395443
## [85]      -Inf      -Inf      -Inf 18.140928 20.168977      -Inf 17.648850
## [92]      -Inf 18.462819 17.607644 18.434298      -Inf 13.035131 18.767153
## [99] 16.873490 16.957907 19.191703 17.051609 16.524105 15.737339 20.592297
## [106] 19.459439 20.291557      -Inf 18.380593 17.750247 18.914626 17.742468
## [113] 19.123569 19.264906 15.175864 18.581842 17.329536 16.569153 17.707331
## [120] 16.692612      -Inf      -Inf 17.238238 15.607868 16.732868      -Inf
```

##	[127]	15.016992	15.261457	14.220976	19.030099	-Inf	-Inf	14.507561
##	[134]	-Inf	16.431574	14.500959	-Inf	19.556308	19.969287	18.543069
##	[141]	16.644609	-Inf	-Inf	18.685802	-Inf	-Inf	17.655956
##	[148]	-Inf	-Inf	17.588113	-Inf	18.168952	16.161473	18.785324
##	[155]	-Inf	19.527489	16.522227	19.013726	18.017033	-Inf	19.426225
##	[162]	18.565674	17.180438	13.223471	-Inf	-Inf	17.549930	-Inf
##	[169]	19.249558	16.927091	15.131314	-Inf	17.848382	18.458857	18.152985
##	[176]	-Inf	-Inf	16.531306	17.148017	-Inf	16.948350	13.558329
##	[183]	-Inf	15.672628	11.549595	17.934373	18.538319	17.348534	16.728399
##	[190]	20.438534	17.373649	17.669375	15.498246	16.184524	-Inf	-Inf
##	[197]	19.146675	19.179268	-Inf	18.836054	-Inf	20.485315	-Inf
##	[204]	-Inf	-Inf	19.476755	16.824939	15.894082	14.172695	-Inf
##	[211]	-Inf	17.801349	16.656768	17.543521	18.623171	18.151922	16.924578
##	[218]	16.786132	17.916960	6.674561	20.334577	14.850869	18.172227	18.902083
##	[225]	13.683466	17.642341	-Inf	16.048084	-Inf	17.036464	18.320953
##	[232]	19.500642	18.328526	14.205981	-Inf	19.224418	-Inf	-Inf
##	[239]	-Inf	18.045687	-Inf	18.938316	17.564718	-Inf	-Inf
##	[246]	18.733782	19.356775	-Inf	-Inf	16.939233	19.764932	18.994959
##	[253]	-Inf	17.072320	-Inf	18.894322	19.357847	-Inf	-Inf
##	[260]	-Inf	16.111081	16.776762	-Inf	2.564949	17.486735	-Inf
##	[267]	18.769750	-Inf	12.142462	19.189232	16.778284	17.534606	-Inf
##	[274]	17.833402	18.386283	19.941717	18.381825	11.461632	17.796825	-Inf
##	[281]	-Inf	16.884527	18.815503	18.048590	18.891175	15.171086	-Inf
##	[288]	18.440483	20.013722	14.865182	18.138252	16.784802	-Inf	19.238243
##	[295]	10.324531	16.647553	17.063518	15.798745	19.202814	-Inf	15.684570
##	[302]	18.464491	14.086051	-Inf	16.182749	17.284398	19.433348	15.988479
##	[309]	-Inf	17.601768	18.236286	-Inf	17.916145	13.649106	18.528106
##	[316]	-Inf	-Inf	-Inf	18.471382	18.545636	-Inf	17.327857
##	[323]	-Inf	-Inf	18.130163	19.725713	18.760467	18.127766	-Inf
##	[330]	-Inf	16.692062	15.333304	18.554171	-Inf	-Inf	17.568478
##	[337]	17.900118	15.740117	15.761421	-Inf	16.199563	16.723348	14.939407
##	[344]	-Inf	20.198577	21.141685	-Inf	19.665512	-Inf	18.908277
##	[351]	15.219111	19.998441	17.118263	-Inf	-Inf	-Inf	17.358224
##	[358]	18.303284	-Inf	19.042155	17.281410	18.425586	18.827945	18.097957
##	[365]	16.972833	19.482103	16.898308	-Inf	14.102366	18.561326	17.113339
##	[372]	16.421984	-Inf	-Inf	17.267238	18.396093	17.984140	-Inf
##	[379]	18.237747	17.822901	17.354058	-Inf	16.131012	18.919042	17.778189
##	[386]	19.752736	-Inf	19.172149	-Inf	-Inf	20.326968	16.140770
##	[393]	17.612071	16.093814	17.434817	-Inf	15.776165	16.486897	19.718422
##	[400]	-Inf	13.816947	17.928234	17.536284	17.133720	19.259871	19.341648
##	[407]	-Inf	-Inf	20.420267	-Inf	17.471278	15.685964	-Inf
##	[414]	-Inf	18.143367	17.978394	20.683485	18.087513	14.711944	18.177063
##	[421]	17.335621	19.945344	-Inf	-Inf	16.597653	17.599700	-Inf
##	[428]	-Inf	-Inf	17.040280	15.201805	4.521789	17.963826	17.515643
##	[435]	18.752518	11.250274	16.352321	17.374703	12.898247	18.819609	17.917730
##	[442]	18.368711	-Inf	-Inf	-Inf	16.185400	14.480078	-Inf
##	[449]	16.428832	-Inf	-Inf	16.828629	17.687294	-Inf	-Inf
##	[456]	19.629347	19.005014	-Inf	18.969072	-Inf	15.087018	-Inf
##	[463]	17.952542	15.725778	19.408088	-Inf	-Inf	14.998337	14.638543
##	[470]	19.158784	-Inf	13.343149	18.139721	16.698581	-Inf	12.310433
##	[477]	18.831576	-Inf	20.132313	16.817093	19.697092	-Inf	-Inf
##	[484]	20.341024	-Inf	17.763960	15.434640	-Inf	16.188109	16.686525
##	[491]	18.974333	15.211206	-Inf	16.554982	-Inf	16.868122	16.367126
##	[498]	-Inf	17.771039	-Inf	16.896812	17.555409	14.862922	-Inf

```
## [505] 14.583088 18.577684 18.574885      -Inf      -Inf      -Inf      -Inf
## [512]      -Inf      -Inf 16.886869 17.466230 15.503670 17.163716      -Inf
## [519] 19.941594      -Inf      -Inf 19.054299 15.285864 12.959280 18.575687
## [526] 15.962027      -Inf      -Inf 17.134872      -Inf      -Inf 14.847931
## [533] 13.367660      -Inf      -Inf 19.568580 13.083230 17.227686      -Inf
## [540] 17.194158 14.697405 18.681755 20.727814 17.453097 16.005673 15.202085
## [547] 17.509453 19.861961      -Inf 20.114340 18.200009 18.874220 10.879669
## [554]      -Inf 19.555111      -Inf 18.427201      -Inf 19.092277 19.187517
## [561] 19.204932 19.028334      -Inf      -Inf      -Inf      -Inf 16.511600
## [568]      -Inf 18.934830      -Inf      -Inf 16.500702 14.862528 16.684899
## [575]      -Inf 16.181305      -Inf 17.716875 13.180723      -Inf 18.755754
## [582] 16.535443      -Inf 18.630195 16.482739 19.314839      -Inf 18.180343
## [589] 16.528317      -Inf 18.606795 20.432730 13.997832      -Inf 19.123908
## [596]      -Inf 19.679466 18.168526 16.712762 18.577519 16.908564      -Inf
## [603]      -Inf 17.281803 20.826329 17.561084      -Inf 19.685508 19.011919
## [610] 18.950350 15.114358 13.302424 18.602509 19.903656 19.479802      -Inf
## [617] 9.768355 18.594637 16.035417 15.761744      -Inf      -Inf 19.953238
## [624]      -Inf 16.769588 17.134973 19.367889 17.517983 11.243502      -Inf
## [631] 15.584971 17.153756 18.821274      -Inf 21.063590 16.611225 17.469526
## [638] 15.802312 16.866026 17.554271      -Inf      -Inf 18.227027      -Inf
## [645] 16.552935 17.606205 12.597270 17.127435      -Inf 17.648949 17.599700
## [652]      -Inf      -Inf 18.627037 16.837359      -Inf 17.846379 15.929235
## [659] 18.868601 16.795819 18.376697      -Inf 16.995562 20.788088      -Inf
## [666] 18.343178      -Inf 18.118984      -Inf 14.877213 20.134456 16.781453
## [673]      -Inf 18.858534 15.560914 11.695247 19.689578      -Inf 15.272255
## [680] 15.561419 20.227356 17.575106 15.950796 17.602130      -Inf      -Inf
## [687]      -Inf 18.420681 20.699460 16.678132 18.664931 16.831857      -Inf
## [694] 16.287856      -Inf 14.768252 16.419816      -Inf      -Inf 17.937777
## [701] 18.055561 16.644650 14.418024 19.142194 19.138519      -Inf 18.242981
## [708] 15.240948      -Inf 17.575092 18.909053 18.676732 19.975086 16.204626
## [715] 17.301219 18.318200 16.173664 15.842011 18.366009      -Inf 16.229321
## [722]      -Inf      -Inf 19.725958      -Inf 17.785802      -Inf      -Inf
## [729] 19.905204      -Inf      -Inf 13.242288 17.632764 15.713643      -Inf
## [736] 17.506887 19.201567      -Inf      -Inf 17.748390 18.259712      -Inf
## [743]      -Inf 18.550459 14.028697      -Inf 17.889581 18.971103 19.630103
## [750] 15.976511 15.051811 18.038881 19.825804      -Inf 12.746479 20.040069
## [757] 18.324290 16.547856 17.035901      -Inf      -Inf 18.777082 18.443971
## [764] 16.317440 16.038779      -Inf 19.193035 13.258308 17.970983      -Inf
## [771]      -Inf 18.739510      -Inf 16.926084 13.311235 19.730475      -Inf
## [778] 19.764217 18.136003      -Inf      -Inf      -Inf 19.273863 16.887634
## [785] 14.631190      -Inf 18.839630      -Inf 17.483451      -Inf 17.048092
## [792]      -Inf      -Inf 17.846178 18.227905      -Inf 19.666086 14.253091
## [799] 19.868466      -Inf 17.750130 18.833580 15.341794 15.016779 16.641206
## [806]      -Inf 15.734606 18.735366
```

```
mean(movies_LogRevenue)
```

```
## [1] -Inf
```

Your Answer:

movies with no revenue will show -inf since the log of 0 does not exist. The mean also shows -inf since the values of -inf are still included when calculating the mean. *Step e*

```

movies_LogRevenueNA <- log(movies4$revenue+1)
movies4$revenue[movies_LogRevenueNA == 0] <- NA
mean(movies_LogRevenueNA, na.rm = TRUE)


```

```
## [1] 11.78363
```

```

hist(movies_LogRevenueNA,
     main = "log of revenue",
     xlab = "Log",
     ylab = "Frequency")

```



Your Answer:

The mean is 17.2798 and the shape of the histogram is left skewed.

Step f

```
skewness(movies4$revenue, na.rm = TRUE)
```

```
## [1] 3.261763
```

```
kurtosis(movies4$revenue, na.rm = TRUE)
```

```
## [1] 16.35398
```

```
skewness(movies_LogRevenueNA, na.rm = TRUE)
```

```
## [1] -0.6475635
```

```
kurtosis(movies_LogRevenueNA, na.rm = TRUE)
```

```
## [1] 1.554936
```

Your Answer:

The skewness is 3.261763 and the kurtosis is 16.35398 for the variable revenue. The skewness is above 0 which explains that it is skewed right. The kurtosis of 16 explains why the histogram has high frequency to the left side and is not flat. The skewness for the log of revenue is -1.946866 which tells us that it is skewed left. The kurtosis of 10.87042 tells us that the that the histogram is not flat.

2 Week 2

1 Is your dataset movies4.tsv the full population, or is it a sample of a larger population? If the latter, how would you describe the full population? **[4 points]**

Your answer here:

The dataset is a sample of a larger population since the index is missing a lot of numbers. We would describe the full population as all movies from 1990 to 2016.

2 For this question, you will assume that your data set is the full population.

- What is the mean of the variable vote_average in your data set? **[2 points]**
- Create a new data set, called movies_sample. Make sure that it is a random sample of your data set of 25 movies. What is the mean of vote_average in this random sample? How does it compare to the mean of vote_average in 2a? **[2 points]**
- In a for loop, create 100 different samples of 25 movies, as in b, and estimate the mean within each sample. Save the mean of each sample in a vector called sample_means. So the first position of the vector would have the mean of the first sample, the second position the mean of the second sample, etc. Print the start of this vector. **[2 points]**
- Summarize and make a histogram of sample_means. What is the mean, standard deviation and shape of its distribution? **[2 points]**
- Using the formulas of the *Central Limit Theorem*, give the distribution of the sample mean of vote_average in a sample of 25 movies from your population (including its expected value and standard deviation). How do these values compare to your answers at d.)? **[2 points]**

Step a

```
#WRITE YOUR CODE HERE  
mean(movies4$vote_average)
```

```
## [1] 6.039356
```

Your Answer:

The mean is 6.039356 *Step b*

```
#WRITE YOUR CODE HERE
movies_sample <- rnorm(25, 6.039356, sd(movies4$vote_average))
mean(movies_sample)
```

```
## [1] 5.888869
```

Your Answer: The mean of the sample set does not match the average of the total movies, because the sample set only includes 25 random movies, and thus is a bit skewed because they are not representative of all the movies.

Step c

```
#WRITE YOUR CODE HERE
sample_means <- c()
for(i in c(1:100)){
  sample_means[i] <- mean(rnorm(25, 6.039356, sd(movies4$vote_average)))}
head(sample_means)
```


```
## [1] 6.000628 6.332331 6.318746 5.969119 5.691468 5.877145
```

Your Answer:

```
5.984406 6.219849 6.091184 6.334625 6.456935 6.357936
```

Step d

```
#WRITE YOUR CODE HERE
hist(sample_means,
      main = "mean of vote average",
      xlab = "amount of movies",
      ylab = "vote average")
```



```
mean(sample_means)
```

```
## [1] 5.979166
```

```
sd(sample_means)
```

```
## [1] 0.2672205
```

Your Answer:

The mean is 6.020662 and the standard deviation is 0.2515497. The shape of the histogram is skewed left.

Step e

Your Answer:

```
m <- mean(movies4$vote_average)
sd <- sd(movies4$vote_average)/sqrt(25)
print(m)
```

```
## [1] 6.039356
```

```
print(sd)
```

```
## [1] 0.2369204
```

We expected a low standard deviation, because of the normal distribution. This is the case because the standard deviation is 0.2515497 so this tells us that the graph in d is indeed normally distributed. The expected value is aswell very close to the middle of the graph also is expected in such a normally distributed graph.

3. Is there a variable in your data set which you would expect to follow a discrete uniform distribution? Explain, without looking at the data, why you expect this data to be discrete uniform. If the data would indeed follow a discrete uniform distribution, what would you expect the mean and standard deviation of this variable to be? Next, plot a histogram of the variable, and plot an appropriate discrete uniform distribution on top of this histogram. Comment whether the variable indeed follows this expected distribution. If the distribution deviates from the discrete uniform distribution, explain why this is the case. **[6 points]**

Your Answer:

I would expect the release month of a movie to follow a discrete uniform distribution because there are 12 possible months, and without looking at the data, each month could be expected to be equally likely for a movie to be released. Therefore, each month would have a probability of $1/12$.

For a discrete uniform distribution from 1 to 12, the expected mean is 6.5 and the expected standard deviation is approximately 3.45. we chose for the release month variable because its logical that the whole year movies were released because there is not a particular month where the benefits are larger to release a new movie. Thats why we expect that release months does not have that much impact.

```
#WRITE YOUR CODE HERE
```


```
summary(movies4$release_month)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   1.000   4.000   7.000   6.708   9.000  12.000
```

```
sd(movies4$release_month)
```

```
## [1] 3.373684
```

```
hist(movies4$release_month,freq=FALSE,
     main="Distribution of Movie release_month",
     xlab="month of release")
abline(h=1/12,col="blue",lwd=2)
```



yes the distribution is similar as the discrete uniform distribution what we would expect.

Your Answer:

Write your formulated response here. There are 808 movies in the dataset, so if the distribution were perfectly uniform, we would expect 67.3 movies to be released in each month. The observed frequencies range from 39 movies in November to 112 movies in September.

The histogram therefore does not follow a perfectly discrete uniform distribution. September has considerably more movie releases than expected, while November has considerably fewer. The deviation from the discrete uniform distribution can be explained by the fact that movie releases are not randomly distributed throughout the year. Movie studios may choose particular release months based on factors such as holidays, audience availability or competition. Therefore, some months can have more movie releases than others.

3 Week 3

For the next part of the assignment, assume that the movies in your data frame are a random sample of a larger population of movies.

1. Give a 90 percent confidence interval for the variable *revenue*. Interpret the result. [2 points]

#WRITE YOUR CODE HERE

```
x <- movies4$revenue
mean(x, na.rm = TRUE)
```

```
## [1] 121314213
```

```
sd(x, na.rm = TRUE)
```

```
## [1] 197480675
```

```
sum(!is.na(x))
```

```
## [1] 551
```

```
t.test(x, conf.level = 0.90)
```

```
##  
## One Sample t-test  
##  
## data: x  
## t = 14.42, df = 550, p-value < 2.2e-16  
## alternative hypothesis: true mean is not equal to 0  
## 90 percent confidence interval:  
## 107452778 135175648  
## sample estimates:  
## mean of x  
## 121314213
```

Your Answer:

The 90% confidence interval for the mean revenue is [107,452,778, 135,175,648].

Assuming that the movies in our data frame are a random sample from a larger population of movies, we are 90% confident that the true mean revenue of the population of movies is between \$107,452,778 and \$135,175,648.

2

- How many movies in your sample were directed by a male director? How many by a female director? **[2 points]**
- Test the null hypothesis that the chance that a movie is being directed by a male or female director is equal. Clearly state your null hypothesis and alternative hypothesis, and your chosen significance level. Use an appropriate R function to test your hypothesis, and state your conclusion. Also give the p-value that corresponds to your test statistic. **[2 points]**
- Recode revenue such that it is measured in millions of dollars. What is the variance of revenue for movies with a male director? What is the variance of revenue for movies with a female director? **[2 points]**
- Assume that the variance of revenue for movies with a male director in your data is the same as the variance for movies with a male director in your population (that is, there is no sampling uncertainty in this estimate). Now, set up a hypothesis test to test for the null hypothesis that the variance of revenue for movies with a female director is equal to the variance for movies with a male director. Clearly state your null hypothesis and alternative hypothesis, and your chosen significance level. Calculate your test statistic and tell how your test statistic is distributed, and state your conclusion. Also give the p-value that corresponds to your test statistic. **[2 points]**

step a

```
#WRITE YOUR CODE HERE
sum(movies4$director_gender == "male", na.rm = TRUE)
```

```
## [1] 685
```

```
sum(movies4$director_gender == "female", na.rm = TRUE)
```

```
## [1] 80
```

Your Answer:

In the sample, 685 movies were directed by a male director, while 80 movies were directed by a female director.

step b

```
#WRITE YOUR CODE HERE
binom.test(685, 765, p = 0.5, alternative = "two.sided")
```

```
##
## Exact binomial test
##
## data: 685 and 765
## number of successes = 685, number of trials = 765, p-value < 2.2e-16
## alternative hypothesis: true probability of success is not equal to 0.5
## 95 percent confidence interval:
## 0.8715412 0.9162048
## sample estimates:
## probability of success
## 0.8954248
```

Your Answer:

The null and alternative hypotheses are:

$H_0: p=0.5$, meaning that the probability of a movie being directed by a male or female director is equal.

$H_A: p \neq 0.5$, meaning that the probabilities are not equal.

I use a significance level of $\alpha=0.05$.

Using a two-sided binomial test in R with `binom.test(685, 765, p = 0.5)`, the p-value is extremely small ($p < 0.001$).

Since the p-value is smaller than the significance level of 0.05, we reject the null hypothesis. There is strong statistical evidence that the probability of a movie being directed by a male director is not equal to the probability of it being directed by a female director.

step c

```
#WRITE YOUR CODE HERE

movies4$revenue_millions <- movies4$revenue / 1000000

var(movies4$revenue_millions[movies4$director_gender == "male"],
    na.rm = TRUE)
```

```
## [1] 39453.26
```

```
var(movies4$revenue_millions[movies4$director_gender == "female"],  
    na.rm = TRUE)
```

```
## [1] 21607.04
```

Your Answer:

Revenue was recoded from dollars to millions of dollars by dividing revenue by 1,000,000.

The variance of revenue for movies with a male director is 39,453.26 (million dollars)².

The variance of revenue for movies with a female director is 21,607.04 (million dollars)².

step d

```
#WRITE YOUR CODE HERE
```

```
F <- 21607.04 / 39453.26  
F
```

```
## [1] 0.5476617
```

```
p_value <- 2 * min(  
  pf(F, df1 = 79, df2 = 79),  
  1 - pf(F, df1 = 79, df2 = 79)  
)  
  
p_value
```

```
## [1] 0.008105479
```

Your Answer:

The null hypothesis is

H₀: standard deviation F₂=39,453.26

meaning that the population variance of revenue for movies with a female director is equal to the population variance for movies with a male director.

The alternative hypothesis is

H_A: standard deviation F₂≠39,453.26.

I use a significance level of α=0.05.

The test statistic is

F=21,607.04 / 39,453.26=0.5477.

Under the null hypothesis, the test statistic follows an F-distribution with 79 degrees of freedom. The two-sided p-value is 0.0081.


Since the p-value (0.0081) is smaller than the significance level of 0.05, we reject the null hypothesis. There is sufficient evidence to conclude that the variance of revenue for movies with female directors is different from the variance of revenue for movies with male directors.

- Create a scatter plot with year on the x axis and mean of revenue within that year on the y-axis. Make sure it has an appropriate title, and appropriate titles and labels for the x- and y-axis. Is revenue of movies increasing or decreasing over time? **[3 points]**
- Recreate the scatter plot with year on the x axis and mean_revenue on the y-axis, but now add bars around each point, indicating the 95% confidence interval. **[3 points]**
- If you look closely at the resulting plot in b.), you will probably see that the confidence intervals in some years are wider than in other years. What are some possible reasons that a confidence interval in a given year can be wide? Can you verify this using your data? **[4 points]**

step a

```
#WRITE YOUR CODE HERE
movies <- read.delim("movies4.tsv", header = TRUE)
mean_revenue <- aggregate(revenue ~ release_year,
                           data = movies,
                           FUN = mean,
                           na.rm = TRUE)

plot(mean_revenue$release_year, mean_revenue$revenue/1000000,
     main = "Mean of revenue",
     xlab = "Year",
     ylab = "Mean Revenue in millions",
     pch = 19)
```



Your Answer:


The overall trend is that revenue of movies is decreasing over time.

step b

```
mean_revenue <- movies %>%
  group_by(release_year) %>%
  summarise(
    mean_revenue = mean(revenue, na.rm = TRUE),
    sd_revenue = sd(revenue, na.rm = TRUE),
    n = sum(!is.na(revenue))
  ) %>%
  mutate(
    se = sd_revenue / sqrt(n),
    lower = mean_revenue - 1.96 * se,
    upper = mean_revenue + 1.96 * se
  )

plot(mean_revenue$release_year,
     mean_revenue$mean_revenue / 1000000,
     main = "Mean Movie Revenue",
     xlab = "Year",
     ylab = "Mean Revenue",
     pch = 19)

arrows(mean_revenue$release_year,
       mean_revenue$lower / 1000000,
       mean_revenue$release_year,
       mean_revenue$upper / 1000000,
       angle = 90,
       code = 3,
       length = 0.05)
```



Your Answer:

The reason for this is that the standard deviations from year to year vary, the higher the standard deviation is the wider the confidence interval needs to be. This is most clearly seen in years with a high standard deviation or a lower N. The N is relevant, because the higher the number of observations you have the narrower the confidence interval is, because you have more data to trust and thus have greater certainty.

step c

```
results <- data.frame(
  year = mean_revenue$release_year,
  sd = mean_revenue$sd_revenue,
  n = mean_revenue$n,
  CI_width = mean_revenue$upper - mean_revenue$lower
)
highest <- results[order(results$CI_width, decreasing = TRUE), ][1:2, ]
lowest <- results[order(results$CI_width), ][1:2, ]

highest
```

```
##   year      sd  n  CI_width
## 5  1994 285619241 11 337580371
## 27 2016 315205264 20 276289595
```

```
lowest
```

```
##      year      sd   n CI_width
## 15 2004 72791859 37 46910279
##  2 1991 30641357  6 49036385
```

4 Week 4

1. There is an argument going on in the movie studio. *Bob* claims that production budgets are getting out of hand, and that the studio should focus on making cheaper movies. *Chantal* disagrees. She tells Bob that “Every dollar we spend on movie production is more than offset by the increase in movie revenue”.
- a. Set up a regression model to test Chantal’s claim, and estimate it. That is, estimate:

$$\text{Revenue}_i = \beta_0 + \beta_1 \text{Budget}_i + \varepsilon_i.$$

Print a summary of your estimated model. **[2 points]**

- b. What is the estimated value of β_1 and how do you interpret it? **[2 points]**
- c. Test for the null hypothesis that $\beta_1 \leq 1$. Report the p-value and state your conclusion. **[2 points]**
- d. Next, estimate the model

$$\text{Log Revenue}_i = \beta_0 + \beta_1 \text{Log Budget}_i + \varepsilon_i.$$

When creating the variables Log Revenue and Log Budget, make sure that movies with a Revenue or Budget of zero are assigned the value “NA”. Print a summary of your estimated model. **[2 points]**

- e. What is the estimated value of β_1 and how do you interpret it? **[2 points]**
- f. Which model has better fit? The level-level model or the log-log model? Explain. **[2 points]**
- g. Who do you think is correct? Bob or Chantal? What would you advise the movie studio to do? **[2 points]**

step a

```
#WRITE YOUR CODE HERE
```

Your Answer:

Write your formulated response here.

step b

```
#WRITE YOUR CODE HERE
```

Your Answer:

Write your formulated response here.

step c

```
#WRITE YOUR CODE HERE
```

Your Answer:

Write your formulated response here.

step d

#WRITE YOUR CODE HERE

Your Answer:

Write your formulated response here.

step e

#WRITE YOUR CODE HERE

Your Answer:

Write your formulated response here.

step f

#WRITE YOUR CODE HERE

Your Answer:

Write your formulated response here.

step g

#WRITE YOUR CODE HERE

Your Answer:

Write your formulated response here.

2

- a. Make a scatterplot with log revenue on the y-axis, and log budget on the x-axis. Within the same scatterplot, draw a straight line which summarizes the association between the log of revenue and log of budget. [3 points]
- b. Using the insights from the scatterplot, your linear model, and your data, find an example of a movie that vastly overperformed. That means, find a movie that had much more revenue than expected, given its budget. For this movie, report its log of revenue, its predicted log of revenue, and its residual. [3 points]

step a

#WRITE YOUR CODE HERE

Your Answer:

Write your formulated response here.

step b

#WRITE YOUR CODE HERE

Your Answer:

Write your formulated response here.

5 Week 5

- Make a frequency table of the variable genre. Which are the three genres that occur the most? **[2 points]**
- Create a new variable called genre2. Make sure that the three most popular genres remain the same, but that all the other genres get categorized into one genre called “Other”. Report a frequency table of this new variable genre2. **[2 points]**
- Estimate an OLS model which has as dependent variable the log of revenue of a movie, and as independent variable the log of budget, a dummy for each level that genre2 can take, and the year of release. Show a summary of the resulting model and interpret each coefficient. **[4 points]**

The movie studio that you work at is releasing a new movie in 2026. It will be a Drama movie with a budget of 20,000,000. It will have a runtime of 120 minutes.

- Estimate a model that is able to predict the revenue of this movie. Give its predicted revenue and include a 90% prediction interval. **[5 points]**
- Plot the residuals against various independent variables in the model. Do you find any evidence of heteroskedasticity? **[2 points]**
- One executive at the studio wants to time the release of the movie to a specific month of the year such that they can maximize revenue. Another executive is of the opinion that the month in which the movie is released will not have a big impact on its revenue. Estimate a model that can test whether the month of release impacts the revenue of a movie. Use an appropriate hypothesis test to test whether month of release matters. What would you advise the movie studio regarding the timing of the release of the movie? **[5 points]**

step a

#WRITE YOUR CODE HERE

Your Answer:

Write your formulated response here.

step b

#WRITE YOUR CODE HERE

Your Answer:

Write your formulated response here.

step c

#WRITE YOUR CODE HERE

Your Answer:

Write your formulated response here.

step d

#WRITE YOUR CODE HERE

Your Answer:

Write your formulated response here.

step e

```
#WRITE YOUR CODE HERE
```

Your Answer:

Write your formulated response here.

step f

```
#WRITE YOUR CODE HERE
```